


Operating path — trained networks only, no branch simulation



trained weights
(training only)

Learning path (offline) — clone state and random stream, run up to three worlds for $H = 3$ h

observed

proposal
reversed

acceptance
reversed

centred Φ
difference
= target